

Animal Classifier

Marko Bukovina a Martin Kvietok

EDA Analysis

DATASET PREVIEW



cane

cane/OIP-rA-sxQehLnbtlaSmqRsaghFj.jpeg

300×225px

brightness=96.9
.jpeg, RGB, 8-bit



cavalo

cavalo/OIP-GtjY8Gnm_4cZwBzbtnP2eAHaE8.jpeg

300×200px

brightness=120.3
.jpeg, RGB, 8-bit



farfalla

farfalla/OIP-HnssZSKnctsXIfiJOgvCWQAHaFi.jpeg

300×225px

brightness=140.1
.jpeg, RGB, 8-bit



gatto

gatto/ec37b10721f01c22d2524518b7444f92e37fe5d404b0144390f8c079a4e5b

6_640.jpg

640×469px

brightness=160.0
.jpg, RGB, 8-bit

```
['cane', 'cavalo', 'elefante', 'farfalla', 'gallina', 'gatto', 'mucca', 'pecora', 'ragno', 'scoiattolo']
```

DATASET OUTLIERS



gatto

gatto/73.jpeg

120×90px

brightness=100.5

.jpeg, RGB, 8-bit

DATASET OUTLIERS



gatto
gatto/73.jpeg
120×90px
brightness=100.5
.jpeg, RGB, 8-bit



farfalla
farfalla/OIP-
3u7uRfPS0Nwr9cUmoRagQwHaEK.jpeg
300×169px
brightness=250.1
.jpeg, RGB, 8-bit



ragno
ragno/eb32b30c2af7003ed1584d05fb1
d4e9fe777ead218ac104497f5c97ca5ecb
5b1_640.jpg
640×479px
brightness=3.0
.jpg, RGB, 8-bit

DATASET OUTLIERS




gatto

gatto/73.jpeg

120×90px

brightness=100.5

.jpeg, RGB, 8-bit



farfalla

farfalla/OIP-

3u7uRfPS0Nwr9cUmoRagQwHaEK.jpeg

300×169px

brightness=250.1

.jpeg, RGB, 8-bit



ragno

ragno/eb32b30c2af7003ed1584d05fb1d4e9fe777ead218ac104497f5c97ca5ecb

5b1_640.jpg

640×479px

brightness=3.0

.jpg, RGB, 8-bit



scioiattolo

scioiattolo/OIP-

117Ca_iCfQgdtUmOUfanvgAAAA.jpeg

127×300px


brightness=85.7

.jpeg, RGB, 8-bit

DATASET OUTLIERS



gatto
gatto/73.jpeg
120×90px
brightness=100.5
.jpeg, RGB, 8-bit



farfalla
farfalla/OIP-
3u7uRfPS0Nwr9cUmoRagQwHaEK.jpeg
300×169px
brightness=250.1
.jpeg, RGB, 8-bit



ragno
ragno/eb32b30c2af7003ed1584d05fb1
d4e9fe777ead218ac104497f5c97ca5ecb
5b1_640.jpg
640×479px
brightness=3.0
.jpg, RGB, 8-bit



scoiattolo
scoiattolo/OIP-
117Ca_iCfogdtUnOUfanvgAAAA.jpeg
127×300px
brightness=85.7
.jpeg, RGB, 8-bit



scimmia

DATASET INFO

Positives

- Large number of samples (**24k**) and an appropriate number of classes (**10**)
- No **corrupted or unreadable images**
- Consistent color models (RGB)
- Correct image formats (.jpeg)
- High data quality for classification tasks
- **Image brightness is not an issue** (outliers are kept in the dataset)

Identified Issues

- Significant **class imbalance**
- 2,672 images with **atypical dimensions** (outliers)
- Minor presence of **incorrectly labeled** samples

Preprocessing


DATA PROCESSING

DATA SPLIT

- 90/10
- BATCHSIZE 64

PREPROCESSING


- 128x128px
- pixel scaling [0,1]
- norm by **mean, std**



DATA AUGMENTATION

- `RandomResizedCrop(target_size)`
- `RandomHorizontalFlip()`
- `RandomRotation(10)`

We also used **minority-augmentation boosting** to balance the dataset (applying stronger augmentation to underrepresented classes)



Model

Configuration

```
EPOCHS = 100

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(img_class_model.parameters(), lr=0.2, momentum=0.9)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, "max", patience=4)
e_stop = EarlyStopping(patience=5, diff=0.01)
```

Configuration

```
EPOCHS = 100

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(img_class_model.parameters(), lr=0.2, momentum=0.9)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, "max", patience=4)
e_stop = EarlyStopping(patience=5, diff=0.01)
```


Model Architecture

```
# Block 1: 256 -> 128
self.block1 = nn.Sequential(
    nn.Conv2d(3, 32, kernel_size=3, padding=1),
    nn.BatchNorm2d(32),
    nn.LeakyReLU(inplace=True),
    nn.Conv2d(32, 32, kernel_size=3, padding=1),
    nn.BatchNorm2d(32),
    nn.LeakyReLU(inplace=True),
    nn.MaxPool2d(2)
)
```



```
# Block 2: 128 -> 64
self.block2 = nn.Sequential(
    nn.Conv2d(32, 64, kernel_size=3, padding=1),
    nn.BatchNorm2d(64),
    nn.LeakyReLU(inplace=True),
    nn.Conv2d(64, 64, kernel_size=3, dilation=2, padding=2),
    nn.BatchNorm2d(64),
    nn.LeakyReLU(inplace=True),
    nn.MaxPool2d(2)
)
```

```
# Block 2: 128 -> 64
self.block2 = nn.Sequential(
    nn.Conv2d(32, 64, kernel_size=3, padding=1),
    nn.BatchNorm2d(64),
    nn.LeakyReLU(inplace=True),
    nn.Conv2d(64, 64, kernel_size=3, dilation=2, padding=2),
    nn.BatchNorm2d(64),
    nn.LeakyReLU(inplace=True),
    nn.MaxPool2d(2)
)
```





```
# Block 3: 64 -> 32
self.block3 = nn.Sequential(
    nn.Conv2d(64, 128, kernel_size=3, padding=1),
    nn.BatchNorm2d(128),
    nn.LeakyReLU(inplace=True),
    nn.Conv2d(128, 128, kernel_size=3, dilation=2, padding=2),
    nn.BatchNorm2d(128),
    nn.LeakyReLU(inplace=True),
    nn.Conv2d(128, 128, kernel_size=3, dilation=3, padding=3),
    nn.BatchNorm2d(128),
    nn.LeakyReLU(inplace=True),
    nn.MaxPool2d(2)
)
```



```
# Block 4: 32 -> 16 (Grad-CAM)
self.block4 = nn.Sequential(
    nn.Conv2d(128, 256, kernel_size=3, padding=1),
    nn.BatchNorm2d(256),
    nn.LeakyReLU(inplace=True),
    nn.Conv2d(256, 256, kernel_size=3, dilation=2, padding=2),
    nn.BatchNorm2d(256),
    nn.LeakyReLU(inplace=True),
    nn.MaxPool2d(2)
)
```



```
# Global average pooling reduces (256×16×16) to (256)
self.global_avg_pool = nn.AdaptiveAvgPool2d((1, 1))

# Classification head
self.fc1 = nn.Linear(256, 128)
self.fc2 = nn.Linear(128, self.numberOfClasses)
```

Training

Beginnings

Image resolution 256x256

```
self.conv1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3, stride=1, padding=1)
self.bn1 = nn.BatchNorm2d(16)

self.conv2 = nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, stride=1, padding=1)
self.bn2 = nn.BatchNorm2d(32)

self.conv3 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1, padding=1)
self.bn3 = nn.BatchNorm2d(64)


self.fc1 = nn.Linear(in_features=64*32*32, out_features=128)
self.fc2 = nn.Linear(in_features=128, out_features=self.numberOfClasses)
```



```
# Block 4: 32 -> 16 (Grad-CAM)
self.block4 = nn.Sequential(
    nn.Conv2d(128, 256, kernel_size=3, padding=1),
    nn.BatchNorm2d(256),
    nn.LeakyReLU(inplace=True),
    nn.MaxPool2d(2)
)
# Global average pooling reduces (256x16x16) to (256)
self.global_avg_pool = nn.AdaptiveAvgPool2d((1, 1))

# Classification head
self.fc1 = nn.Linear(256, 128)
self.fc2 = nn.Linear(128, self.numberOfClasses)

def forward(self, x):
    x = self.block1(x)
    x = self.block2(x)
    x = self.block3(x)
    x = self.block4(x)
    x = self.global_avg_pool(x)
    x = torch.flatten(x, 1)
    x = fun.leaky_relu(self.fc1(x))
    x = self.dropout(x)
    return self.fc2(x)
```




Pred: scoiattolo





151930 fotosearch ©

Train Acc





lyric-energy-11 vague-morning-9
winter-snowball-5 dainty-sunset-2
pretty-star-1



More convolution layers



Final adjustments




Pred: mucca



Pred: scoiattolo



Pred: ragno




Pred: pecora




Mistakes

Pred: gallina



Pred: gallina



Results

KEY FINDINGS

- The model successfully learned to classify animal images with high accuracy (87.6% on the test set)
- Overfitting max 2-3%

